

Tax Evasion and Productivity

Hans Martinez

Western University

September 3, 2026

Recap

Last meeting: What is an interesting question to wrap up the JMP?

- ▶ Counterfactual question: Optimal tax rate

Today:

- ▶ SOC of firm's optimization problem
- ▶ Functional forms for detection probability and tax evasion cost
- ▶ Estimation strategy

Introduction

- ▶ What's the question?
- ▶ Why should you care?

Contribution

- ▶ What do I do?
- ▶ How do I do it?

Overview of Results

- ▶ What do I find?
- ▶ Why is it important?

Colombia

Model

The Model

$$\begin{aligned} \max_{K_{it}, L_{it}, M_{it}, e_{it}} \quad & (1 - q(e_{it})) \left[P_t \mathbb{E}_\varepsilon[Y_{it}] - \rho_t M_{it} - r_t K_{it} - w_t L_{it} \right. \\ & \left. - \tau (P_t \mathbb{E}_\varepsilon[Y_{it}] - \rho_t (M_{it} + e_{it})) \right] \\ & + q(e_{it}) \left[(1 - \tau) (P_t \mathbb{E}_\varepsilon[Y_{it}] - \rho_t M_{it}) - r_t K_{it} - w_t L_{it} \right. \\ & \left. - \phi \rho_t \tau e_{it} \right] - \kappa_{it}(e_{it}, \omega_{it}) \\ \text{s.t.} \quad & Y_{it} = K_{it}^{\alpha_k} L_{it}^{\alpha_L} M_{it}^\beta \exp(\omega_{it} + \varepsilon_{it}) \end{aligned} \tag{1}$$

Important Assumptions

- ▶ Following seminal work of Allingham and Sandmo (1972), tax authority audits firms at random and always detects tax evasion
 - ▶ The probability of detection, $q(e_{it})$, is a function of the level of tax evasion, e_{it} . It might depend on other exogenous variables such as the monitoring technology of the tax authority, for example.
 - ▶ Impose assumption from Empirical model which assumes that choosing M does not enter the probability of detection or the cost of evasion. In other words, choosing M is independent of choosing e.

Important Assumptions (Cont.)

The cost of evasion, $\kappa_{it}(e_{it}, \omega_{it})$: any resources of the firm allocated to the tax evasion activity that could have been employed in a productive activity (e.g. manager's time spent on tax evasion instead of managing the firm).

- ▶ Unlike the penalty of tax evasion, ϕ , it does not depend on whether you are detected or not
- ▶ Increasing in the level of tax evasion, e_{it}
- ▶ Decreasing in firm's productivity, ω_{it} , but spikes at the top of the productivity distribution.

Important Assumptions - discussion on κ

- ▶ More productive firms have more resources to allocate to tax evasion, but since huge firms have access to more sophisticated legal tax avoidance strategies, the cost of illegal tax evasion is higher for the most productive firms.
- ▶ It could also be the case that the probability of detection is flat with respect to productivity except for huge (top productive) firms which could have a probability close to 1.
- ▶ Empirically, without data on cost of evasion or auditing information, q and κ are not separately identified. During estimation, we assume functional forms that allow us to reproduce the empirical patterns of tax evasion across the productivity distribution.

Model (Cont.)

The firms' problem (Equation 1) can be rewritten as follows,

$$\begin{aligned} \max_{K_{it}, L_{it}, M_{it}, e_{it}} \quad & \mathbb{E}[\pi_{it} | \Gamma] = (1 - \tau) (P_t \mathbb{E}[Y_{it}] - \rho_t M_{it}) - r_t K_{it} - w_t L_{it} \\ & + [1 - q(e_{it})] (\rho_t \tau e_{it}) - q(e_{it}) \phi \rho_t \tau e_{it} - \kappa_{it}(e_{it}, \omega_{it}) \\ \text{s.t.} \quad & Y_{it} = K_{it}^{\alpha_K} L_{it}^{\alpha_L} M_{it}^{\beta} \exp(\omega_{it} + \varepsilon_{it}) \end{aligned} \tag{2}$$

First Order Condition (FOC)

An interior solution leads to following first-order conditions,

$$\beta K_{it}^{\alpha_K} L_{it}^{\alpha_L} M_{it}^{-(1-\beta)} \exp(\omega_{it}) = \frac{\rho_t}{P_t} \quad (3)$$

$$\tau \rho_t (1 - (1 + \phi) (q(e_{it}) + q'(e_{it})e_{it})) = \frac{\partial \kappa_{it}(e_{it}, \omega_{it})}{\partial e_{it}} \quad (4)$$

(ε_{ijt} is measurement error, so $\mathcal{E} \equiv E[\exp(\varepsilon_{ijt})] = 1$)

Tax Evasion

To better understand the incentives to overreport inputs, we can rewrite Equation 4 in terms of the evasion elasticity of the probability of detection, $\xi_{q,e} = \frac{\partial q(e_{it})}{\partial e_{it}} \frac{e_{it}}{q(e_{it})}$, as follows,

$$\tau \rho_t \left(1 - (1 + \phi) (1 + \xi_{q,e}) q(e_{it}) \right) = \frac{\partial \kappa_{it}(e_{it}, \omega_{it})}{\partial e_{it}} \quad (5)$$

Corner Solution

Note that some firms will optimally choose not to evade, $e_{it} = 0$.

For these firms, the marginal cost of tax evasion is greater than the marginal benefit for any non-zero level of tax evasion.

This implies that for these firms, Equation 4 becomes,

$$\tau \rho_t \leq \frac{\partial \kappa_{it}(0, \omega_{it})}{\partial e_{it}} \quad (6)$$

Second Order Condition (SOC)

To guarantee a unique solution, the hessian matrix of the firm's problem must be negative definite.

$$\nabla^2 \mathbb{E}[\pi_{it} | \Gamma] = \begin{bmatrix} \frac{\partial^2 \mathbb{E}[\pi_{it} | \Gamma]}{\partial e_{it}^2} & 0 & 0 & 0 \\ 0 & \frac{\partial^2 \mathbb{E}[\pi_{it} | \Gamma]}{\partial M_{it}^2} & \frac{\partial^2 \mathbb{E}[\pi_{it} | \Gamma]}{\partial M_{it} \partial L_{it}} & \frac{\partial^2 \mathbb{E}[\pi_{it} | \Gamma]}{\partial M_{it} \partial K_{it}} \\ 0 & \frac{\partial^2 \mathbb{E}[\pi_{it} | \Gamma]}{\partial L_{it} \partial M_{it}} & \frac{\partial^2 \mathbb{E}[\pi_{it} | \Gamma]}{\partial L_{it}^2} & \frac{\partial^2 \mathbb{E}[\pi_{it} | \Gamma]}{\partial L_{it} \partial K_{it}} \\ 0 & \frac{\partial^2 \mathbb{E}[\pi_{it} | \Gamma]}{\partial K_{it} \partial M_{it}} & \frac{\partial^2 \mathbb{E}[\pi_{it} | \Gamma]}{\partial K_{it} \partial L_{it}} & \frac{\partial^2 \mathbb{E}[\pi_{it} | \Gamma]}{\partial K_{it}^2} \end{bmatrix} \quad (7)$$

- ▶ For a negative definite hessian matrix, the leading principal minors must alternate signs,
 - ▶ odd order minors must be negative and even-order minors positive.
- ▶ In our case, since
 - ▶ 1) the production function is strictly concave (CD) and
 - ▶ 2) choosing the level of tax evasion does not depend on the production function inputs; second cross derivatives are zero,
 - ▶ we would only have to check the first leading principal minor —the second derivative with respect to e_{it} — is negative,

$$\frac{\partial^2 \mathbb{E}[\pi_{it} | \Gamma]}{\partial e_{it}^2} = -\tau \rho_t (1 + \phi) (2q'(e_{it}) + e_{it} q''(e_{it})) - \frac{\partial^2 \kappa_{it}(e_{it}, \omega_{it})}{\partial e_{it}^2} < 0 \quad (8)$$

For strict concavity, we need only one of the following two conditions to be strictly positive,

$$\frac{\partial^2 \kappa_{it}(e_{it}, \omega_{it})}{\partial e_{it}^2} \geq 0 \quad (9)$$

$$\tau \rho_t (1 + \phi) (2q'(e_{it}) + e_{it} q''(e_{it})) \geq 0 \quad (10)$$

The second condition can be simplified to

$$-\frac{q''(e_{it})}{q'(e_{it})}e_{it} < 2 \quad (11)$$

since economic theory implies that $q'(e_{it}) > 0$ and $e_{it} > 0$, this condition imposes a **restriction on the curvature** of the detection function, $q(e_{it})$

- ▶ **Linear** detection **always** satisfies the SOC, since $q''(e_{it}) = 0$.
- ▶ **Convex** also satisfies the SOC **always**, since $q''(e_{it}) > 0$.
- ▶ **Concave** works if not too concave, i.e. **inequality holds**.

What does this inequality imply in practice?

With $\phi = 0$, logs of the interior FOC $\tau\rho_t(1 - q - q'e) = \partial\kappa/\partial e$ define

- ▶ the **log marginal benefit** $h(e) \equiv \ln(\tau\rho_t) + \ln(1 - q(e) - q'(e)e)$,
- ▶ with FOC $h(e) = \ln(\partial\kappa/\partial e)$.

	$q(e)$	SOC	$h(e) - \ln(\tau\rho_t)$	FOC ceiling
Linear	$\min\{\lambda e, 1\}$	$0 < 2$	$\ln(1 - 2\lambda e)$	$\lambda e < \frac{1}{2}$
Concave	$1 - e^{-\lambda e}$	$\lambda e < 2$	$-\lambda e + \ln(1 - \lambda e)$	$\lambda e < 1$
Convex	$\min\{(\lambda e)^k, 1\}$ $k > 1$	$1 - k < 2$	$\ln(1 - (1+k)(\lambda e)^k)$	$\lambda e < (1+k)^{-1/k}$

- ▶ All satisfy the SOC on their interior (exp's $\lambda e < 2 \iff \lambda e < 1$);
- ▶ All impose a restriction on λe .

A detection technology: the audit threshold

- ▶ The firm chooses overreporting e .
- ▶ The tax authority draws a threshold $T \sim U[0, 1/\lambda]$ and audits the firm if $e \geq T$.
 - ▶ An audit always uncovers evasion (Allingham and Sandmo 1972).
 - ▶ T is random to the firm (audit selection, third-party cross-checks, which inspector).

$$q(e) = \Pr(T \leq e) = \min\{\lambda e, 1\}.$$

- ▶ **Threshold scale:** $1/\lambda$ is the largest overreporting any audit draw tolerates; the *average* audit-triggering threshold is $1/2\lambda$.
- ▶ **Evasion bound:** the FOC stops firms at $\lambda e < \frac{1}{2}$
 - ▶ the most aggressive evader overreports below the *average* audit-triggering threshold,
 - ▶ at threshold, $q(1/2\lambda) = 50\%$ (a coin flip)

If $q(e)$ is concave, $T \sim \text{Exp}(\lambda)$ and $q(e) = 1 - e^{-\lambda e}$, mean threshold $1/\lambda$, detection $\approx 63\%$ at $e = 1/\lambda$

How likely is the most aggressive evader caught?

At the FOC ceiling, the largest evasion drives the marginal benefit to zero,
 $q(b) + q'(b)b = 1$.

	b	$q(b)$	$\xi_{q,e} _{e=b}$	$E[T]$
Linear	$\frac{1}{2\lambda}$	50%	1 (unit elastic)	$\frac{1}{2\lambda}$
Exponential	$\frac{1}{\lambda}$	63%	$\frac{\lambda b}{e^{\lambda b} - 1} < 1$ (inelastic)	$\frac{1}{\lambda}$
Power	$\frac{1}{\lambda(1+k)^{1/k}}$	$\frac{1}{1+k} < 50\%$	$k > 1$ (elastic)	$\frac{k}{\lambda(k+1)}$

- ▶ $\xi_{q,e} = q'(e)e/q(e)$ is the evasion elasticity of the detection probability
- ▶ $\frac{1}{e-1} = 0.58$

Discussion

- ▶ Power (convex) form produces the lowest average audit-triggering threshold,
 - ▶ Exponential (concave) the largest.
 - ▶ Linear is intuitive: at max evasion, the tax authority flips a coin to audit/caught the firm
- ▶ An increase in 1% in the level of evasion increases the probability of being caught, by
 - ▶ 1% if linear (constant)
 - ▶ more than 1% if convex (constant)
 - ▶ less than 1% if concave (increasing), $\xi_{q,e}|_{e=0} = 0$ and $\xi_{q,e}|_{e=1/\lambda} = 0.58$.
- ▶ The extra dollar the firm hides increase its probability of detection more than the previous dollar.
- ▶ For linear and exponential, the FOC bounds evasion *at* the expected value of T
 - ▶ For power, FOC bounds evasion *below* its expected value: At $k = 2$, $\frac{1}{\lambda\sqrt{3}} < \frac{2}{\lambda 3}$

Estimation and recommendation

- ▶ **Linear** and **exponential** carry a single parameter λ
- ▶ **Power** $(\lambda e)^k$ adds a parameter, curvature k

Do **linear** and **power**,

- ▶ Start with **linear** ($k = 1$). Linear is the $k = 1$ case of the power form (ceiling $(1+k)^{-1/k} = \frac{1}{2}$).
- ▶ Then, estimate power (λ, k) , imposing feasibility $\lambda e < (1+k)^{-1/k}$ for *every* firm. Test if $k = 1$.

Empirical Strategies

- ▶ Testing for the presence of tax evasion by industry
 - ▶ Robust to different specifications
- ▶ Regression
 - ▶ Useful to study tax evasion heterogeneity across firms and over time
- ▶ Deconvolution
 - ▶ Recover distribution of tax evasion and productivity
 - ▶ Estimation of unbiased Production Function

Data

Testing

Empirical Strategy

Functional Forms

The probability of detection, $q(e_{it})$, and the cost of evasion, $\kappa_{it}(e_{it}, \omega_{it})$, functions should display the following properties to be consistent with the observed patterns in the data and with theory:

- ▶ $q(e_{it}) = \min\{\lambda e, 1\}$
 - ▶ $q(e_{it}) \in [0, 1]$, $q(0) = 0$, $q(\infty) = 1$, and $q(e_{it})$ is increasing in e . Since I do not have variation in government audit effort, λ is treated as constant across firms and time
- ▶ $\kappa_{it}(e_{it}, \omega_{it}) = e_{it} \mathbf{e}^{\delta_0 - \delta_1 \omega_{it} + \delta_2 \omega_{it}^2 + \psi_{it}}$
 - ▶ $\kappa_{it}(e_{it}, \omega_{it}) \in [0, \infty)$, $\kappa_{it}(0, \omega_{it}) = 0$, $\frac{\partial \kappa_{it}(0, \omega_{it})}{\partial e_{it}} \geq 0$.
 - ▶ $\frac{\partial \kappa_{it}(e_{it}, \omega_{it})}{\partial \omega_{it}} \leq 0$ and $\frac{\partial^2 \kappa_{it}(e_{it}, \omega_{it})}{\partial \omega_{it}^2} \geq 0$ to be consistent with the observed patterns in the data.
 - ▶ U-shape not ideal, but simple and captures the idea of decreasing cost at high ω
 - ▶ Corporations out of the FOC
 - ▶ ψ_{it} idiosyncratic shock to the cost of evasion to explain tax evasion heterogeneity observed in empirical analysis

FOC

Using the functional forms defined above, the FOCs wrt tax evasion become for unincorporated firms,

$$\tau\rho_t(1 - 2\lambda e_{it}) = \mathbf{e}^{\delta_0 - \delta_1\omega_{it} + \delta_2\omega_{it}^2 + \psi_{it}}$$

Taking logs and defining

$$h(e_{it}) \equiv \ln(\tau\rho_t) + \ln(1 - 2\lambda e_{it}), \quad e_{it} \in (0, \frac{1}{2\lambda}),$$

the evasion FOC become

$$\psi_{it} = h(e_{it}) - \delta_0 + \delta_1\omega_{it} - \delta_2\omega_{it}^2 \tag{12}$$

How to estimate?

- ▶ $\tau\rho_t$ is *computed from data*

I do not observe e_{it} nor ω_{it} at the firm level.

- ▶ However, using the estimated **parameters of the production** function, we can
- ▶ rewrite the FOC in terms of data, parameters, and
- ▶ *only one* independent unobservable ε

How to estimate the production function parameters and the distribution of ε ? **Same as before**

Estimating β

Same as before: From Corporations, the truth-reporters, I recover

- ▶ the output elasticity of materials, β , and the
- ▶ distribution of the error term, f_ε ,

For Corporations, $M_{it}^* = M_{it}$, and thus,

$$\ln \left(\frac{\rho_t M_{it}^*}{P_t Y_{it}} \right) = \ln(\beta) - \varepsilon_{it}$$

Production Function Estimation

- Having recovered β , for everyone, I can form the following variable,

$$\begin{aligned}\mathcal{V}_{it} &\equiv \ln \left(\frac{\rho_t M_{it}^*}{P_t Y_{it}} \right) - \ln(\beta) \\ &= \ln \left(\frac{\rho_t M_{it}^*}{P_t Y_{it}} \right) - \ln \left(\frac{\rho_t M_{it}}{P_t Y_{it}} \right) + \ln \left(\frac{\rho_t M_{it}}{P_t Y_{it}} \right) - \ln(\beta) \\ &= \ln \left(\frac{M_{it}^*}{M_{it}} \right) - \varepsilon_{it}\end{aligned}\tag{13}$$

Production Function Estimation (Cont.)

- To recover the remaining production function parameters, the **original strategy is still valid**

To see why, note that,

$$\begin{aligned}\mathcal{W}_{it} &\equiv y_{it} - \beta(m_{it}^* - \mathcal{V}_{it}) \\ &= y_{it} - \beta(m_{it}^* - \ln\left(\frac{M_{it}^*}{M_{it}}\right) + \varepsilon_{it}) \\ &= y_{it} - \beta(m_{it}^* - \ln(M_{it}^*) + \ln(M_{it})) - \beta\varepsilon_{it} \\ &= y_{it} - \beta m_{it} - \beta\varepsilon_{it} \\ &= \alpha_K k_{it} + \alpha_L l_{it} + \omega_{it} + (1 - \beta)\varepsilon_{it}\end{aligned}\tag{14}$$

the last equality is *exactly the same* as the original discussed strategy

Using productivity's Markov process to recover remaining PF parameters

Therefore, the **previous approach is still valid** to recover the remaining production function parameters, productivity's Markov process, and the distribution of the productivity shock $(\alpha_K, \alpha_L, \omega_{it})$.

Assuming, productivity follows an AR(1) Markov process,

$$\omega_{it} = \gamma_0 + \gamma_1 \omega_{it-1} + \eta_{it} \quad (15)$$

we recover the PF parameters, by substituting Equation 14, into Equation 15,

$$\begin{aligned} \mathcal{W}_{it} = & \gamma_0 + \alpha_K k_{it} + \alpha_L l_{it} \\ & + \gamma_1 (\mathcal{W}_{it-1} - \alpha_K k_{it-1} - \alpha_L l_{it-1}) \\ & + (1 - \beta) \varepsilon_{it} - \gamma_1 (1 - \beta) \varepsilon_{it-1} + \eta_{it} \end{aligned} \quad (16)$$

Estimating Equation 16

By finding instruments, we can estimate Equation 16, using the conditional moment,

$$E[(1 - \beta)\varepsilon_{it} - \gamma_1(1 - \beta)\varepsilon_{it-1} + \eta_{it} | Z_{it}] = 0$$

- I use capital and labour as instruments for themselves and reported materials and $\mathcal{W}_{it-2}$ as instruments to control for endogeneity between $\mathcal{W}_{it-1}$ and

$$\varepsilon_{it-1}$$

- $Z = \{k_{it}, l_{it}, m_{it-1}^*, \mathcal{W}_{it-2}\}$

Using the estimated parameters of the production function

As previously discussed, with α_K and α_L at hand form,

$$\tilde{\mathcal{W}}_{it} \equiv \mathcal{W}_{it} - \alpha_K k_{it} - \alpha_L l_{it} = \omega_{it} + (1 - \beta)\varepsilon_{it} \quad (17)$$

and recover f_ω using deconvolution,

$$f_{\tilde{\mathcal{W}}} = \int f_\varepsilon \left(\frac{1}{1 - \beta} (\tilde{\mathcal{W}} - \omega) \right) f_\omega(\omega) d\omega$$

Solving for M and e

NEW From Equation 13, I can *algebraically solve* for true inputs as function of observed variables and an the error term,

► $M_{it} = M_{it}^* \mathbf{e}^{-\mathcal{V}_{it} - \varepsilon_{it}}.$

Therefore,

$$\begin{aligned} e_{it} &= M_{it}^* - M_{it} \\ &= M_{it}^* (1 - \mathbf{e}^{-\mathcal{V}_{it} - \varepsilon_{it}}) \end{aligned}$$

- ε_{it} is not observed firm-by-firm, only its distribution f_ε
- not a literal point recovery of e_{it} , but the map that could be evaluated for some draw of $\varepsilon^{(s)} \sim \hat{f}_\varepsilon$

Recast in terms of one unobservable

I can do the same and solve for ω_{it} from Equation 17

For each value of ε everything else is pinned down:

$$e_{it}(\varepsilon) = M_{it}^* (1 - \mathbf{e}^{-\mathcal{V}_{it} - \varepsilon}) \quad (18)$$

$$\omega_{it}(\varepsilon) = \tilde{\mathcal{W}}_{it} - (1 - \beta)\varepsilon \quad (19)$$

$$\psi_{it}(\varepsilon) = h(e_{it}(\varepsilon)) - \delta_0 + \delta_1 \omega_{it}(\varepsilon) - \delta_2 \omega_{it}(\varepsilon)^2. \quad (20)$$

MSL

- ▶ We could estimate f_ψ and $\theta = \{\lambda, \delta_0, \delta_1, \delta_2\}$ by Maximum Simulated Likelihood (MSL)
- ▶ After the recast, one unobservable remains: ε_{it} , with known distribution f_ε (from stage 1),
- ▶ we also know the distribution of ω
- ▶ I would need a *density transformation* to get from the primitives' density $f_{\psi, \omega, \varepsilon}(\psi, \omega, \varepsilon)$ to the density of the observables $f_{\mathcal{V}, \tilde{\mathcal{W}}}(\mathcal{V}, \tilde{\mathcal{W}})$
 - ▶ accounting for the nonlinear map from ψ to $\mathcal{V}$ (Theorem 7.5.3, next slide) *and*
 - ▶ integrating out ε <!--

$$f_{\mathcal{V}_{it}, \tilde{\mathcal{W}}_{it}}(\mathcal{V}_{it}, \tilde{\mathcal{W}}_{it}) = \int f_\psi(\psi(\varepsilon_{it})) f_\omega(\omega(\varepsilon_{it})) f_\varepsilon(\varepsilon_{it}) |\det J\Phi^{-1}(\varepsilon)| d\varepsilon_{it}$$

The object to estimate is f_ψ and $\theta = (\lambda, \delta_0, \delta_1, \delta_2)$, since f_ω and f_ε are already known →

Theorem 7.5.3 Change of Variables in R^n Judd (1998)

Assume that $\Phi : X \subset R^n \rightarrow R^n$ is C^1 and that there is a C^1 inverse map $\Phi^{-1} : Y \subset R^n \rightarrow X$ where $Y = \Phi(X)$ and Y and X are open. If f is integrable on X , then

$$\int_Y f(y)dy = \int_X f(\Phi(x)) |\det J(\Phi)(x)| dx,$$

where $J(\Phi)(x)$ is the Jacobian of Φ at x .

Corollary: Density transformation

Let $f_X(x)$ be the joint pdf of X , B any region in the support of Y , and $A = \Phi^{-1}(B)$

For $Y = \Phi(X)$: $P(X \in A) = P(\Phi(X) \in \Phi(A)) = P(Y \in B)$. Apply the theorem with $\Psi \equiv \Phi^{-1}$ in place of Φ (Φ^{-1} is C^1 with C^1 inverse Φ) to $g = f_X \cdot \mathbf{1}_A$:

$$P(X \in A) = \int_A f_X(x) dx = \int_B f_X(\Phi^{-1}(y)) |\det J\Phi^{-1}(y)| dy.$$

Since B is arbitrary, and using $J\Phi^{-1}(y) = [J\Phi(x)]^{-1}$ (inverse function theorem) at $x = \Phi^{-1}(y)$, then, the joint pdf of Y is:

$$f_Y(y) = \frac{f_X(x)}{|\det J\Phi(x)|} \bigg|_{x=\Phi^{-1}(y)} \quad (21)$$

Target Density

To find the correct target density of observables, note that, given (K_{it}, L_{it}) , the variables $(\mathcal{V}_{it}, \tilde{\mathcal{W}}_{it})$ form an affine map in the (Y_{it}, M_{it}^*) -space,

$$\mathcal{V}_{it} = m_{it}^* - y_{it} + c_1 \quad (13 \text{ rearranged})$$

$$\tilde{\mathcal{W}}_{it} = y_{it} - \beta(m_{it}^* - \mathcal{V}_{it}) - \alpha_K k_{it} - \alpha_L l_{it} \quad (\text{equivalent to 17})$$

where $c_1 \equiv \ln \left(\frac{\rho_t}{P_t \beta} \right)$

Substituting the first equation into the second cancels the m_{it}^* terms exactly:

$$\tilde{\mathcal{W}}_{it} = (1 - \beta) y_{it} + \beta c_1 - \alpha_K k_{it} - \alpha_L l_{it} \quad (22)$$

Φ is a linear bijection in (y_{it}, m_{it}^*) -space

$$\Phi : (y_{it}, m_{it}^*) \mapsto (\mathcal{V}_{it}, \tilde{\mathcal{W}}_{it})$$

$$\begin{pmatrix} \mathcal{V} \\ \tilde{\mathcal{W}} \end{pmatrix} = \underbrace{\begin{pmatrix} -1 & 1 \\ 1 - \beta & 0 \end{pmatrix}}_A \begin{pmatrix} y \\ m^* \end{pmatrix} + \begin{pmatrix} c_1 \\ \beta c_1 - \alpha_K k - \alpha_L l \end{pmatrix}.$$

Since $\det A = -(1 - \beta) \neq 0$ ($0 < \beta < 1$) $\implies \Phi$ is an affine bijection

- ▶ $\implies (\mathcal{V}_{it}, \tilde{\mathcal{W}}_{it})$ contain the same information as (Y_{it}, M_{it}^*)
- ▶ Therefore, the correct target density $f_{\mathcal{V}, \tilde{\mathcal{W}}|K, L}$
- ▶ **No need** to include M^* in the joint density or condition on $(\mathcal{V}$ contains all information of $M^*)$

Conditioning on ω_{it-1} instead of (K_{it}, L_{it})

Due to the Markov process of productivity, $\omega_{it} = g(\omega_{it-1}) + \eta_{it}$,
 $\eta_{it} \perp\!\!\!\perp (\omega_{it-1}, k_{it}, l_{it})$,

$$f_{\omega_{it}|\omega_{it-1}, K_{it}, L_{it}}(\omega_{it}) = f_{\omega_{it}|\omega_{it-1}}(\omega_{it}) = f_{\eta}(\omega_{it} - g(\omega_{it-1}))$$

Conditional on ω_{it-1} , the only randomness left in ω_{it} is η_{it} , which doesn't depend on (K_{it}, L_{it})

► ω_{it-1} is a **sufficient statistic**: once it's known, (K_{it}, L_{it}) adds nothing.

Setting Φ in the $(\psi_{it}, \omega_{it}, \varepsilon_{it})$ -space

$\Phi : (\psi_{it}, \omega_{it}, \varepsilon_{it}) \mapsto (\mathcal{V}_{it}, \tilde{\mathcal{W}}_{it}, a)$, padding with the auxiliary $a \equiv \varepsilon_{it}$ so
 $\Phi : \mathbb{R}^3 \rightarrow \mathbb{R}^3$ matches the theorem's dimension requirement:

Define $G(\psi_{it}, \omega_{it}) \equiv u(e^*(\psi_{it}, \omega_{it}))$,

- ▶ where $u(e) = \ln((M_{it} + e)/M_{it})$, holding M_{it} fixed as e varies $(M \perp e \mid \omega)$, and
- ▶ $e^*(\psi_{it}, \omega_{it}) = h^{-1}(\delta_0 - \delta_1 \omega_{it} + \delta_2 \omega_{it}^2 + \psi_{it})$ inverts the FOC for e_{it} ,

$$\mathcal{V} = G(\psi, \omega) - \varepsilon, \quad \tilde{\mathcal{W}} = \omega + (1 - \beta)\varepsilon, \quad a = \varepsilon$$

- ▶ Like (K_{it}, L_{it}) before, ω_{it-1} does not enter the map

Conditional density

Let $J_i(\varepsilon_{it}) = |\det J\Phi^{-1}|$, then,

$$f_{\mathcal{V}_{it}, \tilde{\mathcal{W}}_{it} | \omega_{it-1}}(\mathcal{V}_{it}, \tilde{\mathcal{W}}_{it}) = \int f_{\psi}(\psi(\varepsilon_{it})) J_i(\varepsilon_{it}) f_{\eta}(\omega(\varepsilon_{it}) - g(\omega_{it-1})) f_{\varepsilon}(\varepsilon_{it}) d\varepsilon_{it}. \quad (23)$$

Assuming $\psi_{it}, \varepsilon_{it} \perp (K_{it}, L_{it}) \mid \omega_{it-1}$ ## Since ω_{it-1} isn't observed, use $\tilde{\mathcal{W}}_{it-1}$

$\tilde{\mathcal{W}}_{it-1}$ is firm- i *data*

- ▶ ε_{it-1} unobserved, but with known density f_{ε} (same as stage 1, by stationarity) and
- ▶ already assumed independent $\varepsilon_{it-1} \perp \varepsilon_{it}, \eta_{it}$ in stage 1,
- ▶ extending $\varepsilon_{it-1} \perp \psi_{it}$
- ▶ therefore, we can marginalize ε_{it-1} with an ordinary integral,

Recovering f_η

Define the residual used in the AR(1) productivity process,

$$\zeta_{it} \equiv \mathcal{W}_{it} - \alpha_K k_{it} - \alpha_L l_{it} - \gamma_0 - \gamma_1 (\mathcal{W}_{it-1} - \alpha_K k_{it-1} - \alpha_L l_{it-1}) = (1 - \beta) \varepsilon_{it} - \gamma_1$$

a sum of three mutually independent variables, two with known densities, f_ε .
Fourier deconvolution recovers the third:

$$\phi_\eta(t) = \frac{\phi_\zeta(t)}{\phi_\varepsilon((1 - \beta)t) \phi_\varepsilon(-\gamma_1(1 - \beta)t)} \longrightarrow f_\eta,$$

The Jacobian matrix

$$J\Phi = \frac{\partial(\mathcal{V}, \tilde{\mathcal{W}}, a)}{\partial(\psi, \omega, \varepsilon)} = \begin{pmatrix} G_\psi & G_\omega & -1 \\ 0 & 1 & 1 - \beta \\ 0 & 0 & 1 \end{pmatrix}$$

where $G_\psi \equiv \partial G / \partial \psi$ and $G_\omega \equiv \partial G / \partial \omega$

Determinant

Expand along column 1 (only G_ψ is nonzero there):

$$\det J\Phi = G_\psi \cdot \det \begin{pmatrix} 1 & 1 & -\beta \\ 0 & & 1 \end{pmatrix} = G_\psi \cdot 1 = G_\psi$$

- ▶ G_ω never enters because every cofactor multiplying it comes from a zero in column 1.
- ▶ Mechanically: $\tilde{\mathcal{W}} = \omega + (1 - \beta)\varepsilon$ has **unit coefficient on ω** , a flat translation, so it contributes no curvature.
- ▶ A Jacobian only appears where the map is genuinely curved: the $\psi \rightarrow \mathcal{V}$ link through G , not $\omega \rightarrow \tilde{\mathcal{W}}$.

Unpacking G_ψ

$G = u \circ e^*$, so by the chain rule $G_\psi = u'(e^*) \cdot e_\psi^*$.

► $u'(e)$: from $u(e) = \ln(M_{it} + e) - \ln M_{it}$, with M_{it} fixed,

$$u'(e) = \frac{1}{M_{it} + e}$$

evaluated at $e = e^*$, $M_{it} + e^* = M_{it}^*$, $\implies u'(e) = 1/M_{it}^*$

► e_ψ^* : differentiate the FOC identity $\psi = h(e^*(\psi, \omega)) - \delta_0 + \delta_1\omega - \delta_2\omega^2$ w.r.t. ψ (implicit function theorem, valid since $h' \neq 0$):

$$1 = h'(e^*) \cdot e_\psi^* \implies e_\psi^* = \frac{1}{h'(e^*)}$$

Combine:

$$G_\psi = \frac{1}{M_{it}^* h'(e)}$$

Assembling $J_i(\varepsilon)$

From Equation 21, the multiplying factor is $1/|\det J\Phi| = 1/|G_\psi|$:

$$J_i(\varepsilon) \equiv \frac{1}{|G_\psi|} = |h'(e)| M_{it}^*$$

Every piece is traceable: - $|h'(e)|$ from inverting the FOC ($\psi \rightarrow e$), - M_{it}^* from the log-ratio definition ($e \rightarrow \mathcal{V}$, holding the *true* M_{it} fixed) - they multiply because $G = u \circ e^*$ is a **composition**, so the chain rule stacks them.

Marginalizing out the auxiliary

Applying the density formula gives the joint density of $(\mathcal{V}, \tilde{\mathcal{W}}, a)$, but only $(\mathcal{V}, \tilde{\mathcal{W}})$ are observed;

- ▶ Integrating a out (an ordinary single integral, not part of the theorem) gives the likelihood contribution of firm i :

$$f_{\mathcal{V}, \tilde{\mathcal{W}}}(\mathcal{V}_{it}, \tilde{\mathcal{W}}_{it}) = \int f_{\psi}(\psi(\varepsilon)) J_i(\varepsilon) f_{\omega}(\omega(\varepsilon)) f_{\varepsilon}(\varepsilon) d\varepsilon \quad (25)$$

A 3-for-3 change of variables that ends in a 1-D integral:

- ▶ the theorem handles the square transformation and hands back a Jacobian;
- ▶ marginalizing the unobserved auxiliary coordinate is the separate, ordinary step left over once the theorem's job is done.

Simulated log-likelihood

$$\mathcal{L}(\theta, f_\psi) = \sum_{i \in \text{Unincorp}} \ln \left(\int f_\psi(\psi(\varepsilon)) J_i(\varepsilon) f_\omega(\omega(\varepsilon)) f_\varepsilon(\varepsilon) d\varepsilon \right) \quad (26)$$

Simulating the likelihood

Approximate Equation 25 by **Monte Carlo**

$$\widehat{\mathcal{L}}_i(\theta) = \frac{1}{S} \sum_{s=1}^S f_{\psi}(\psi(\varepsilon^{(s)})) J_i(\varepsilon^{(s)}) f_{\omega}(\omega(\varepsilon^{(s)})), \quad \varepsilon^{(s)} \stackrel{iid}{\sim} \hat{f}_{\varepsilon} \quad (27)$$

- ▶ Draws $\varepsilon^{(s)}$ come from *exactly* the density $\hat{f}_{\varepsilon}(\varepsilon) d\varepsilon$ multiplying the integrand — this is plain Monte Carlo integration, unbiased for **any** $S \geq 1$
- ▶ Alternatively, use **deterministic quadrature** to integrate Equation 25

Why this works despite ε entering e nonlinearly

- ▶ $e(\varepsilon) = M_{it}^* (1 - \mathbf{e}^{-\mathcal{V}_{it} - \varepsilon})$ is nonlinear in ε
- ▶ But Equation 27 averages the **density value** $f_{\psi}(\psi(\varepsilon^{(s)})) J_i(\varepsilon^{(s)}) f_{\omega}(\omega(\varepsilon^{(s)}))$ at each draw
 - ▶ never averages $\psi(\varepsilon^{(s)})$ *by itself*
- ▶ Monte Carlo integration of a (possibly wild) function against its own sampling density is unbiased regardless of the function's shape
 - ▶ nonlinearity only breaks things if the reweighting is skipped
- ▶ Every simulated $\varepsilon^{(s)} \neq$ the firm's true ε_{it} , almost surely
 - ▶ but draws implying an implausible (ψ, ω) pair get down-weighted by $f_{\psi} \cdot J_i \cdot f_{\omega}$; only draws near the truth carry weight,
 - ▶ so the weighted average reconstructs the correct posterior-weighted integral even though no single draw is the truth

Convergence — a known theorem

Lee (1992): if the simulator of a likelihood contribution is (i) unbiased and continuous in θ , and (ii) built from draws i.i.d. across firms and replicates, then under standard M-estimation regularity (compactness, identification, dominance) the simulated-MLE is consistent and $\sqrt{N}$ -equivalent to the infeasible exact MLE, **provided** $\sqrt{N}/S \rightarrow 0$ as $N, S \rightarrow \infty$ (fixed S leaves an $O(1/S)$ bias, since $\ln(\text{average}) \neq \text{average of } \ln$)

- ▶ Unbiasedness (i): shown above, holds by construction
- ▶ Continuity (ii): h , q , and $J_i(\varepsilon)$ are smooth throughout the interior domain $e \in (0, \frac{1}{2\lambda})$
- ▶ Draws i.i.d. across i, t, s , independent of θ : by construction

An even cleaner argument: quadrature

- ▶ The preferred implementation is **deterministic quadrature**, not Monte Carlo simulation, for this 1-D integral
- ▶ Quadrature avoids simulation noise because it is deterministic and can be driven arbitrarily small by adding nodes
- ▶ The $\sqrt{N}/S \rightarrow 0$ condition above is then moot: with negligible quadrature error, ordinary MLE consistency applies directly

Caveat: this is a two-step estimator

- ▶ β and f_ε are **first-stage estimates**, not the truth
- ▶ Still consistent under standard conditions, but naive stage-2 standard errors ignore stage-1 estimation uncertainty
- ▶ Use the Murphy and Topel (1985) two-step correction, or bootstrap both stages jointly

Would GMM work ?

Substituting and simplifying

$$\begin{aligned}\psi_{it} + \delta_1(1 - \beta)\varepsilon = & \ln(\tau\rho_t) + \ln(1 - 2\lambda M_{it}^*(1 - \mathbf{e}^{-\mathcal{V}_{it} - \varepsilon})) \\ & - \delta_0 + 2\delta_2(1 - \beta)\tilde{\mathcal{W}}_{it}\varepsilon - \delta_2(1 - \beta)^2\varepsilon^2 \\ & + \delta_1\tilde{\mathcal{W}}_{it} - \delta_2\tilde{\mathcal{W}}_{it}^2\end{aligned}\tag{28}$$

Taking conditional expectation

Using $\psi \perp Z$, $\varepsilon \perp Z$, and $\omega \perp \varepsilon$ (the same independence MSL already relies on).
 $\varepsilon \perp Z \Rightarrow E[\delta_1(1 - \beta)\varepsilon \mid Z] = \delta_1(1 - \beta)E[\varepsilon \mid Z] = 0$

$$\begin{aligned} 0 = & \ln(\tau\rho_t) + E[\ln(1 - 2\lambda M_{it}^*(1 - \mathbf{e}^{-\mathcal{V}_{it} - \varepsilon})) \mid Z] \\ & \underbrace{- \delta_0 - \delta_2(1 - \beta)^2\sigma_\varepsilon^2}_{\text{Constant (known from stage 1)}} \\ & + \delta_1\tilde{\mathcal{W}}_{it} - \delta_2\tilde{\mathcal{W}}_{it}^2 \end{aligned} \tag{29}$$

Why not GMM with instruments?

Everything in Equation 29 is a known function of data and θ , **except**:

$$\ell_{it}(\varepsilon; \lambda) \equiv \ln(1 - 2\lambda M_{it}^* (1 - \mathbf{e}^{-\mathcal{V}_{it} - \varepsilon}))$$

so the moment reads

$$0 = \ln(\tau \rho_t) + E[\ell_{it}(\varepsilon; \lambda) \mid Z] - \delta_0 - \delta_2(1 - \beta)^2 \sigma_\varepsilon^2 + \delta_1 \tilde{\mathcal{W}}_{it} - \delta_2 \tilde{\mathcal{W}}_{it}^2$$

Simulating the nonlinear term

If we hold firm it 's own $(M_{it}^*, \mathcal{V}_{it})$ fixed and simulate ε ,

$$\hat{\ell}_{it}^S(\lambda) = \frac{1}{S} \sum_{s=1}^S \ell_{it}(\varepsilon^{(s)}; \lambda), \quad \varepsilon^{(s)} \stackrel{iid}{\sim} \hat{f}_{\varepsilon}$$

By the LLN, as $S \rightarrow \infty$:

$$\hat{\ell}_{it}^S(\lambda) \xrightarrow{a.s.} \bar{\ell}_{it}(\lambda) \equiv \int \ell_{it}(\varepsilon; \lambda) f_{\varepsilon}(\varepsilon) d\varepsilon$$

A specific, computable number given $(M_{it}^*, \mathcal{V}_{it}, \lambda)$. Looks like it solved the problem — no more unobserved ε .

But it's the wrong number

The moment needs $\ell_{it}(\varepsilon_{it}; \lambda)$

— firm it 's value at its **own realized** ε_{it} — not the average over every value ε could have taken.

$\ell_{it}(\cdot; \lambda)$ is nonlinear in ε , so generically

$$\bar{\ell}_{it}(\lambda) \neq \ell_{it}(\varepsilon_{it}; \lambda)$$

Let $\eta_{it}(\lambda) \equiv \ell_{it}(\varepsilon_{it}; \lambda) - \bar{\ell}_{it}(\lambda) \neq 0$. Swapping in $\bar{\ell}_{it}(\lambda)$ silently injects $-\eta_{it}(\lambda)$ into every firm's moment.

Conclusion: this doesn't work

- ▶ $\eta_{it}(\lambda)$ is a nonlinear function of firm it 's own ε_{it} , entangled with $(M_{it}^*, \mathcal{V}_{it})$
- ▶ For Z to be *relevant* for λ , it must correlate with the firm's evasion behavior $(M_{it}^*, \mathcal{V}_{it})$, but $\eta_{it}(\lambda)$ depends on them
- ▶ $\varepsilon \perp Z$ alone doesn't deliver $E[Z \cdot \eta_{it}(\lambda_0)] = 0$
- ▶ Nonlinearity plus averaging over the unconditional marginal throws away exactly the firm-specific information tying ε_{it} to its own data

Recasting in M

The same trap reappears if we instead simulate $M_{it}^{(s)} = \beta P_t Y_{it} / \rho_t \cdot \mathbf{e}^{-\varepsilon^{(s)}}$ and form $e_{it}^{(s)} = M_{it}^* - M_{it}^{(s)}$ — this looks like it avoids $\mathcal{V}_{it}$, but

$$e(\varepsilon) = M_{it}^* (1 - \mathbf{e}^{-\mathcal{V}_{it} - \varepsilon}) = M_{it}^* - \underbrace{\frac{\beta P_t Y_{it}}{\rho_t} \mathbf{e}^{-\varepsilon}}_{= M(\varepsilon)} = M_{it}^* - M(\varepsilon)$$

using $\mathbf{e}^{-\mathcal{V}_{it}} = \beta P_t Y_{it} / (\rho_t M_{it}^*)$. **It's the identical map** — $\mathcal{V}_{it}$'s information is just M_{it}^* and $P_t Y_{it} / \rho_t$, so the M -version reconstructs it implicitly. Same nonlinearity, same missing reweighting, same problem.

What a valid simulated-moments design would need

- ▶ Match a **data-computable, ψ -free** population object against a forward-simulated analog — e.g. $E[\mathcal{V}_{it}|Z] = E[u_{it}|Z]$ is directly computable from data (since $\varepsilon \perp Z$), while $u^{(s)} = u(e^*(\psi^{(s)}, \omega^{(s)}; \theta))$ can be simulated forward from trial $(\psi^{(s)}, \omega^{(s)})$ draws
- ▶ This ties back to the deconvolved f_u from stage 1 as the natural target
- ▶ Delivers θ but not f_ψ , and needs its own machinery — no real gain over quadrature-MLE, which is already a cheap 1-D integral
- ▶ **Not pursued further** — MSL/quadrature is the headline strategy

An Alternative: ELVIS (Schennach (2014))

- ▶ The “Why not GMM” failure wasn’t fundamental — it was a *naive-averaging* bug: swapping in $\bar{\ell}_{it}(\lambda)$ silently drops the firm-specific reweighting
- ▶ Schennach (2014, *Econometrica*, “Entropic Latent Variable Integration via Simulation”) targets exactly this class: $E[g(U, Z, \theta)] = 0$ with U unobserved, g arbitrary — nonlinear, nonmonotone, nonseparable
- ▶ Her own canonical example (errors-in-variables: unobserved true regressor X^* , observed $X = X^* + \Delta X$) is structurally our M vs. M^*

Recasting the moment in M

Take M_{it} (true materials) — not ε_{it} — as the one unobservable. Everything else is a deterministic function of M and data:

$$e(M) = M_{it}^* - M, \quad \varepsilon(M) = \ln(M_{it}^*/M) - \mathcal{V}_{it}, \quad \omega(M) = \tilde{\mathcal{W}}_{it} - (1 - \beta)\varepsilon(M)$$

$$g(M, Z_{it}, \theta) \equiv \psi(M) = h(e(M)) - \delta_0 + \delta_1 \omega(M) - \delta_2 \omega(M)^2 \quad (30)$$

- ▶ Support $\mathcal{U}_{Z, \theta} = (M_{it}^* - \frac{1}{2\lambda}, M_{it}^*)$ — the feasibility ceiling $e < 1/(2\lambda)$ becomes the domain, not a complication (ELVIS explicitly allows θ -dependent support)
- ▶ Same ψ as before — recasting in M only changes *which variable ELVIS treats as unobserved*, not the moment itself

The entropic fix

Schennach's Theorem 2.1: the infinite-dimensional question “does *some* distribution of M given Z satisfy $E[g] = 0$ ” is **equivalent** to a finite-dimensional system in (θ, γ) :

$$\inf_{\mu} \|E_{\mu \times \pi}[g(M, Z, \theta)]\| = 0 \iff \inf_{\gamma \in \mathbb{R}^{d_g}} \|E_{\pi}[\tilde{g}(Z, \theta, \gamma)]\| = 0$$

$$\tilde{g}(Z, \theta, \gamma) \equiv \frac{\int g(M, Z, \theta) \exp(\gamma' g(M, Z, \theta)) d\rho(M|Z)}{\int \exp(\gamma' g(M, Z, \theta)) d\rho(M|Z)} \quad (31)$$

- ▶ The exponential tilt $\exp(\gamma' g(\cdot))$ **is** the missing reweighting — not assumed, but *solved for* as a Lagrange multiplier on the moment constraint
- ▶ $\rho(M|Z)$ is any user-chosen dominating measure — by her Remark 2.3, its choice has **zero effect** on the resulting estimate, only support matters

Estimation

Simulate $M^{(s)} \sim \rho(\cdot | Z_{it})$ (Metropolis) and form a ratio of averages:

$$\hat{g}(Z_{it}, \theta, \gamma) = \frac{\frac{1}{S} \sum_{s=1}^S g(M^{(s)}, Z_{it}, \theta) \exp(\gamma' g(M^{(s)}, Z_{it}, \theta))}{\frac{1}{S} \sum_{s=1}^S \exp(\gamma' g(M^{(s)}, Z_{it}, \theta))} \quad (32)$$

- ▶ Solve $(\hat{\theta}, \hat{\gamma})$ jointly by standard GMM on $E[\hat{g}(Z_{it}, \theta, \gamma)] = 0$
- ▶ λ identified the same way as before: sector–time variation in $\tau \rho_t$ shifting the evasion FOC's benefit
- ▶ 15 moments (finalized 2026-08-28), all built from independence assumptions the model already maintains — $\theta_{\text{smooth}} = (\delta_0, \delta_1, \delta_2, \sigma_\psi^2)$ (a new nuisance parameter, ψ 's own dispersion), λ gridded separately; full list in `Code/Rcpp/1200-stage2-elvis.cpp` / `CLAUDE.md`

ELVIS vs. MSL vs. plain MSM

	recovers f_ψ ?	needs f_ε ?	valid despite nonlinearity?
MSL (headline)	yes	yes	yes (by construction)
plain MSM	no	yes	no — naive-averaging bug
ELVIS	no	no — only support of M	yes (Thm 2.1)

- ▶ ELVIS's real payoff: stage 1's deconvolution of f_ε becomes **optional** for stage 2 — only $\hat{\beta}$ is needed to construct $\mathcal{V}, \tilde{\mathcal{W}}$
- ▶ $\hat{f}_\varepsilon$ can still be reused to build an efficient ρ (less simulation noise) — a computational choice, not an identification requirement

Going Further: One Unified Estimation

- ▶ If ELVIS works for stage 2 alone, it works for the *whole model* at once
- ▶ **There is only one true unobservable:** M_{it} — true materials, for every firm-period, corp and unincorporated alike
- ▶ No more $\mathcal{V}$, $\tilde{\mathcal{W}}$, no Jacobian, no density-transformation theorem — none of that machinery is needed once M itself is the latent
- ▶ Also sidesteps what made MSL “messier”: the AR(1) correction there needed the *conditional density* $f_{\omega_{it}|\omega_{it-1}}$ (i.e. f_{η} , recovered by Fourier deconvolution). ELVIS only needs **moment orthogonality** $E[\eta_{it}Z_{it}] = 0$ — no density of η at all

From M to ε

The materials FOC pins ε down exactly, given M :

$$\varepsilon_{it}(M_{it}) = \ln \beta - \ln \left(\frac{\rho_t M_{it}}{P_t Y_{it}} \right) \quad (33)$$

From M to ω — GNR's own trick, made exact

Substituting Equation 33 into the production identity

$\omega_{it} = y_{it} - \alpha_K k_{it} - \alpha_L l_{it} - \beta \ln M_{it} - \varepsilon_{it}$, the y_{it} terms **cancel exactly**:

$$\omega_{it}(M_{it}) = (1 - \beta) \ln M_{it} - \alpha_K k_{it} - \alpha_L l_{it} - \ln \beta + \ln \rho_t - \ln P_t \quad (34)$$

- ▶ This is **not information loss** — it's Gandhi-Navarro-Rivers's own $\Phi_{it} = y_{it} - \hat{\varepsilon}_{it}$ step, specialized to constant (Cobb-Douglas) elasticity: their Φ loses y_{it} the *same* way in this special case
- ▶ Standard proxy-variable logic (Olley-Pakes/Levinsohn-Petrin/ACF/GNR): under the timing assumption, optimal m is a monotonic function of ω *alone* — so m already perfectly proxies ω . y_{it} 's only job is revealing ε_{it} (via the share equation); once that's done, y_{it} has nothing further to say about ω
- ▶ y_{it} is used exactly **once** — inside $\varepsilon_{it}(M_{it})$

The AR(1) moment

Raw form (before substituting the FOC) — note the correct coefficients:

$$y_{it} - \beta \ln M_{it} - \alpha_K k_{it} - \alpha_L l_{it} - \gamma_0 - \gamma_1 (y_{it-1} - \beta \ln M_{it-1} - \alpha_K k_{it-1} - \alpha_L l_{it-1}) = \varepsilon_{it}$$

Since ε is itself just a function of M (Equation 33), substitute Equation 34 directly — no ε left anywhere:

$$\eta_{it}(M_{it}, M_{it-1}) = \omega_{it}(M_{it}) - \gamma_0 - \gamma_1 \omega_{it-1}(M_{it-1}) \quad (35)$$

Evasion moments, by firm type

Unincorporated (interior FOC), same as before:

$$e_{it}(M_{it}) = M_{it}^* - M_{it}, \quad \psi_{it}(M_{it}) = h(e_{it}(M_{it})) - \delta_0 + \delta_1 \omega_{it}(M_{it}) - \delta_2 \omega_{it}(M_{it})$$

Corp — relaxed from a hard constraint to a centering moment:

$$E[e_{it}(M_{it}) \mid \text{Corp}_i] = 0 \quad \text{or} \quad \text{Med}[e_{it}(M_{it}) \mid \text{Corp}_i] = 0$$

- ▶ No longer “corp never evade” imposed *outside* the estimation — corp can evade a little, just centered at zero, matching what the data already show (near-zero, imprecise evasion at both tails)

Unincorporated — evade on average, as an inequality:

$$E[e_{it}(M_{it}) \cdot \mathbb{1}(\text{Unincorp}_i) - U_2] = 0, \quad U_2 \geq 0$$

- ▶ ELVIS’s own device for inequalities (Example 1.3 in Schennach (2014)) — no separate technique needed

The full moment system

Moment	Object	Identifies	Instruments
Materials FOC AR(1) innovation	$\varepsilon_{it}(M_{it})$ $\eta_{it}(M_{it}, M_{it-1})$	β (mainly via corp) $\alpha_K, \alpha_L, \gamma_0, \gamma_1$	constant / sector k_{it}, l_{it} , lagged m^*
Evasion FOC (unincorp)	$\psi_{it}(M_{it})$	$\lambda, \delta_0, \delta_1, \delta_2$	sector-time $\times \tau \rho_t$
Corp evasion	$e_{it}(M_{it})$	pins $M_{it} \approx M_{it}^*$ for corp	— (centering)
Unincorp evasion sign	$e_{it}(M_{it}) \cdot$ $\mathbb{1}(\text{Unincorp})$	tightens the identified set	— (inequality)

- ▶ All firm-periods, corp and unincorporated, enter **jointly** — no data thrown away at any stage
- ▶ One draw of M_{it} (and its lag) feeds the *whole* moment vector — not double-counting, this is standard for any GMM system with multiple moments per observation

More moments from assumptions we already make

Every independence assumption already maintained — but never yet written as a moment — adds an (almost free) row to g :

Assumption	New moment
$\varepsilon \perp \psi$	$E[\varepsilon_{it}(M_{it})\psi_{it}(M_{it})] = 0$
$\omega \perp \varepsilon$	$E[\omega_{it}(M_{it})\varepsilon_{it}(M_{it})] = 0$
$\psi \perp \omega$	$E[\psi_{it}(M_{it})\omega_{it}(M_{it})] = 0$ — overidentifying check on the U-shape
$\varepsilon_{it} \perp \varepsilon_{it-1}$	$E[\varepsilon_{it}(M_{it})\varepsilon_{it-1}(M_{it-1})] = 0$ — free, M_{it-1} already drawn
$\varepsilon_{it-1} \perp \psi_{it}$	$E[\varepsilon_{it-1}(M_{it-1})\psi_{it}(M_{it})] = 0$ — free, same draw

- Formal license: Schennach (2014)'s own supplement, Prop. B.1 (conditional-mean-zero $\square$ countable unconditional moments $E[Y e^{i\xi X}] = 0$) and Prop. B.2 (independence, not just uncorrelatedness, factors the same way 77 / 137

What about the counterfactual?

MSL's revenue integral weights by $f_{\psi}(\psi(\varepsilon))J_i(\varepsilon)f_{\omega}(\omega(\varepsilon))f_{\varepsilon}(\varepsilon)$ — needs the full recovered density, which ELVIS doesn't deliver.

- ▶ ELVIS's own machinery for exactly this situation: Schennach's device for bounding a functional of the unobservable that isn't already one of the estimating moments — her worked example is literally recovering $E[U]$ this way (Schennach (2014), p.1250)
- ▶ Reusing the tilted density from the *main* estimation as if it were “the” posterior would give a number, but understate uncertainty — that tilted distribution is only the *entropy-maximizing* member of the set of distributions consistent with the moments, not necessarily the true one (she is explicit about this). The honest version re-solves for the functional's own identified set

Implementing it: an auxiliary moment

Hold $\hat{\theta}$ fixed (already estimated — the counterfactual is a forward simulation, not new estimation). Reuse the closed-form counterfactual evasion FOC already established:

$$\tilde{e}_i(\tilde{\tau}) = \tilde{e}(\tilde{\tau}; \omega_i(M_i), \psi_i(M_i)) = \frac{\tilde{\tau}\rho_t - C(\omega_i, \psi_i)}{2\lambda\tilde{\tau}\rho_t}, \quad C(\omega, \psi) \equiv \mathbf{e}^{\delta_0 - \delta_1\omega + \delta_2\omega^2}$$

(true M_i unchanged under the counterfactual — materials FOC doesn't depend on τ , per the model's own maintained assumptions) and form firm i 's counterfactual revenue as a function of the *same* latent M_i :

$$R_i(\tilde{\tau}; M_i) = \tilde{\tau} \left[P_t Y_{it} - \rho_t M_i - (1 - q(\tilde{e}_i(\tilde{\tau}))) \rho_t \tilde{e}_i(\tilde{\tau}) \right]$$

For a candidate $\tilde{\tau}$, introduce an auxiliary parameter $\tilde{\theta}(\tilde{\tau}) \equiv \frac{1}{N} \sum_i E[R_i(\tilde{\tau}) \mid Z_i]$ and **append its own moment** to the system:

$$E[R_i(\tilde{\tau}; M_i) - \tilde{\theta}] = 0$$

with θ fixed at $\hat{\theta}$ — only $\tilde{\theta}$ and an extended α (one more component, for the new

Point or set — decided by the procedure, not assumed

For a fixed $\tilde{\tau}$, grid over **candidate revenue values** $\tilde{\theta}$ (not over θ this time):

1. For each candidate $\tilde{\theta}$, solve the inner problem — maximize the objective (EL, or the CUE-type quadratic $-\frac{1}{2}\hat{g}'\hat{V}^{-1}\hat{g}$) over the extended γ , with the $R_i(\tilde{\tau}; M_i) - \tilde{\theta}$ moment included
2. That maximum is a criterion function $\hat{L}_n(\tilde{\theta}; \tilde{\tau})$ — exactly the same construction used for θ itself
3. **Profile against the unrestricted model:**
 $\hat{Q}_n(\tilde{\theta}; \tilde{\tau}) = -(\hat{L}_n(\tilde{\theta}; \tilde{\tau}) - \sup_{\theta', \gamma} \hat{L}_n(\theta', \gamma))$ — subtracting the global max, not using $\hat{L}_n$ raw, is what gives the efficiency gain this construction is built for
4. Region: $\{\tilde{\theta} : n\hat{Q}_n(\tilde{\theta}; \tilde{\tau}) \leq c_\alpha\}$. c_α via subsampling (Chernozhukov–Hong–Tamer 2007) for a tight region, or a fast conservative first pass: $c_\alpha = \chi_{d_g, \alpha}^2$ (table lookup, no subsampling — skips step 3's profiling)

► Repeat across a grid of $\tilde{\tau}$: **the shape decides**. A sharp trough at every $\tilde{\tau} \rightarrow$ 80 / 137

Implicit Assumptions

- ▶ The functional forms for $q(e)$ and $\kappa(e)$ are correctly specified and do not change over time
- ▶ No entry, no exit
- ▶ Corporations never evade
- ▶ Partial equilibrium, Inelastic demand, (No GE effects)
 - ▶ Tax increases do not change optimal input choices and output. FOC of materials does not depend on tax rate
 - ▶ Firms are already assumed to be price takers

What about τ ?

- ▶ The data includes sales taxes paid on purchases, $\tau \rho_t M_{it}^*$ and sales taxes paid on sales, $\tau P_t Y_{it}$.
- ▶ I can estimate at the firm level the sales tax rate on sales $\tau_{S,it} = \frac{\tau P_t Y_{it}}{P_t Y_{it}}$ and the sales tax rate on purchases $\tau_{P,it} = \frac{\tau \rho_t M_{it}^*}{\rho_t M_{it}^*}$.
- ▶ I can also estimate the reported effective sales tax rate $\tau_{R,it} = \frac{\tau P_t Y_{it} - \tau \rho_t M_{it}^*}{P_t Y_{it}} = \tau_{S,it} - \tau_{P,it} \frac{\rho_t M_{it}^*}{P_t Y_{it}}$
- ▶ and the *corrected* effective sales tax rate $\tau_{C,it} = \tau_{S,it} - \tau_{P,it} \frac{\rho_t M_{it}^*}{P_t Y_{it}} = \tau_{S,it} - \tau_{P,it} \beta e^{-\varepsilon_{it}}$
 - ▶ I use in regressions $\tau_{\tilde{C},it} = \tau_{S,it} - \tau_{P,it} \beta$

Is the variation in the different estimated tax rates relevant

- ▶ I use the following regression model

$$\mathcal{V}_{it} = \gamma_j + \gamma_t + \tau_{X,it} + \zeta_{it}$$

- ▶ where γ_j are industry fixed effects,
- ▶ γ_t are year fixed effects, and
- ▶ $\tau_{X,it}$ is one of the estimated tax rates ($\tau_{S,it}$, $\tau_{P,it}$, $\tau_{R,it}$, or $\tau_{\tilde{C},it}$).

Variation in estimated τ predicts variation in tax evasion across firms

Dependent Variable: Model:	(1)	(2)	$\mathcal{V}$ (3)	(4)	(5)
<i>Variables</i>					
τ_S	-0.2804* (0.1333)		0.5053*** (0.1422)		
τ_P		-1.890*** (0.1252)	-2.026*** (0.1258)		
τ_R				-1.387* (0.6846)	
τ_C					0.9046*** (0.1803)
<i>Fixed-effects</i>					
Industry (3-digit)	Yes	Yes	Yes	Yes	Yes
Year	Yes	Yes	Yes	Yes	Yes
<i>Fit statistics</i>					
Observations	40,490	40,505	40,490	40,490	40,490
R ²	0.04767	0.09016	0.09215	0.07638	0.05335
Within R ²	0.00068	0.04532	0.04736	0.03081	0.00665

Clustered (plant & Year) standard-errors in parentheses

*Signif. Codes: ***: 0.01, **: 0.05, *: 0.1*

Couterfactual Question

What Would Be The Optimal Tax Rate?

The government faces a trade-off:

- ▶ as the government increases the tax rate, the tax revenue increases
- ▶ as the tax rate increases, the marginal benefit of evading taxes increase leading to higher levels of tax evasion which reduces tax revenue

What is the optimal rate that maximizes expected tax revenue for the government?

What is the tax rate that lead to median (media) tax evasion of zero?

The Government's Problem

The government wants to choose the sales tax rate τ that maximizes the expected tax revenue from the sum of all firm i

$$\max_{\tau} \sum_i q(e) \tau [P_t Y_{ijt} - \rho_t M_{ijt}] + (1 - q(e)) \tau [P_t Y_{ijt} - \rho_t (M_{ijt} + e_{ijt})]$$

Simplifying,

$$\max_{\tau} \sum_i \tau [P_t Y_{ijt} - \rho_t M_{ijt} - (1 - q(e)) \rho_t e_{ijt}]$$

where the first two terms are the potential tax revenue (in the absence of tax evasion) and the last terms is the expected undetected tax evasion.

for firm i ,

$$\begin{aligned} & [P_t Y_{ijt} - \rho_t M_{ijt} - (1 - q(e)) \rho_t e_{ijt}] \\ & + \tau \left[-(1 - q(e)) \rho_t \frac{\partial e_{ijt}}{\partial \tau} + \rho_t e_{ijt} \frac{\partial q(e)}{\partial e} \frac{\partial e_{ijt}}{\partial \tau} \right] = 0 \end{aligned}$$

simplifying,

$$\begin{aligned} P_t Y_{ijt} - \rho_t M_{ijt} &= (1 - q(e)) \rho_t e_{ijt} \\ &+ \tau \rho_t \frac{\partial e_{ijt}}{\partial \tau} \left[(1 - q(e)) - \frac{\partial q(e)}{\partial e} e_{ijt} \right] \end{aligned}$$

Numerical Strategy

1. Having recovered $\theta = (\lambda, \delta_0, \delta_1, \delta_2)$ and f_ψ through MSL are held fixed for everything below
2. For each firm-year, recover the implied **type** (ω, ψ) at the *observed* τ , using the same map as the likelihood (Equation 25): for $M_{it}(\varepsilon) = M_{it}^* \mathbf{e}^{-\mathcal{V}_{it}-\varepsilon}$,

$$\omega(\varepsilon) = \tilde{\mathcal{W}}_{it} - (1 - \beta)\varepsilon, \quad \psi(\varepsilon) = h(e(\varepsilon)) - \delta_0 + \delta_1\omega(\varepsilon) - \delta_2\omega(\varepsilon)^2$$

3. Hold (ω, ψ) fixed as firm primitives (they don't respond to the tax rate) and resolve the evasion FOC at a proposed $\tilde{\tau}$ to find the counterfactual evasion, $\tilde{e}$. Under linear q , closed form, no root-finding:

$$\tilde{e}(\tilde{\tau}; \omega, \psi) = \frac{\tilde{\tau} \rho_t - C(\omega, \psi)}{2\lambda \tilde{\tau} \rho_t}, \quad C(\omega, \psi) \equiv \mathbf{e}^{\delta_0 - \delta_1 \omega + \delta_2 \omega^2 + \psi}$$

4. $\tilde{e}$ is a **deterministic function of the same** ε that generated (ω, ψ) . So expected revenue is a *single* integral over ε , weighted exactly as the likelihood weights it, normalized by the firm's own likelihood contribution $\mathcal{L}_i(\theta)$ (already computed in stage 2), the conditional expectation given firm i 's data:

$$E[R_i(\tilde{\tau})] = \frac{1}{\mathcal{L}_i(\theta)} \int \tilde{\tau} \left[P_t Y_{it} - \rho_t M_{it}(\varepsilon) - (1 - q(\tilde{e}(\varepsilon))) \rho_t \tilde{e}(\varepsilon) \right] A(\varepsilon) d\varepsilon \quad (36)$$

where $A = f_\psi(\psi(\varepsilon)) J_i(\varepsilon) f_\omega(\omega(\varepsilon)) f_\varepsilon(\varepsilon)$

5. Evaluate Equation 36 with the same quadrature nodes used in stage 2
6. Find $\tilde{\tau}^* = \arg \max_{\tilde{\tau}} \sum_i E[R_i(\tilde{\tau})]$ by grid or golden-section search over the scalar $\tilde{\tau}$

Stage-2 ELVIS: Preliminary Results

Implementation

Moment set A (row list authoritative in `Code/Rcpp/1200-stage2-elvis.cpp`'s header comment): $\theta_{\text{smooth}} = (\lambda, \delta_0, \delta_1, \delta_2, \eta)$, 8 moments ($d_g = 8$), no nuisance parameters (additional to γ 's):

$$g(M, Z, \theta_{\text{smooth}}) = \begin{pmatrix} \psi(M, Z) \\ \varepsilon(M, Z) \\ \psi(M, Z) \ln M \\ \psi(M, Z) \omega(M, Z) \\ \psi(M, Z) \omega(M, Z)^2 \\ \varepsilon(M, Z) \ln M \\ \varepsilon(M, Z) e(M, Z) \\ \varepsilon(M, Z) \omega(M, Z) \end{pmatrix}$$

Implementation Details

- ▶ **Objective:** CUE quadratic, $\hat{L}_n(\theta, \gamma) = \frac{1}{2} \hat{g}' \hat{\Omega}^- \hat{g}$, $\hat{\Omega}^-$ an eigenvalue-truncated generalized inverse (drops near-zero eigenvalues of the tilted-moment covariance) — same construction AK2020 use
- ▶ **Sampler:** per-firm Metropolis chain (RcppParallel, n_burn=500, n_keep=1000), deterministic per-firm RNG reseed — same (θ, γ) always sees the same chain
- ▶ **Solver:** BOBYQA, unbounded γ (Schennach 2014, p.359 — a large/degenerate γ is an expected solution, not a numerical failure), two-pass refinement (warm-started from pass 1)


Implementation

- ▶ **λ search, two phases:** (i) a quantile grid, $\lambda = p/\text{quantile}(M^*, 0.5)$ for $p \in \{0.01, \dots, 0.9\} \cap (0, 0.5]$; (ii) an automatic linear-refinement grid bracketing the quantile grid's minimum, warm-started from the fit *at* that minimum
- ▶ **Sample:** 32,378 unincorporated firm-periods (3,340 at the $\tau_P = 0$ corner, included as centered non-evaders), two productivity-instrument choices for $(\alpha_K, \alpha_L, \beta)$

$\hat{L}_n(\lambda)$ — Quantile Grid + Linear Refinement

Stage-2 ELVIS: CUE objective across the lambda grid

Moment set A, n_burn=500, n_keep=1000, corner_mode=include_zero



Reading the picture:

- ▶ $\mathcal{W}_{it-2}$'s minimum sits below the smallest λ its refinement grid reached — real evidence it needs to extend further.
- ▶ m_{it-1}^* is different: its quantile-phase minimum was *interior* ($\lambda = 2.351 \times 10^{-6}$), so the refinement window was built from immediate neighbors, not an edge extension — it never actually probed below 1.175×10^{-6} at all. Its own minimum sits at that window's edge, but nothing smaller has been tested, so whether it too needs to extend is an open question, not yet evidence either way.
- ▶ A further refinement pass — genuinely new territory for m_{it-1}^* , a continued extension for $\mathcal{W}_{it-2}$ — is the natural next step before treating either $\hat{\lambda}$ as final.

Headline Estimates

Instru- ment	$\hat{\lambda}$	$\hat{L}_n$	$\hat{\delta}_0$	$\hat{\delta}_1$	$\hat{\delta}_2$	$\hat{\omega}^* = \hat{\delta}_1 / 2\hat{\delta}_2$	$\max \gamma $
m_{it-1}^*	1.18×10^{-6}	0.0108	-1.099	1.453	0.218	3.34	2.08
$\mathcal{W}_{it-2}$	5.88×10^{-7}	0.0073	-1.305	1.859	0.396	2.35	1.43

$\hat{\delta}_2 > 0$ for both — the cost function opens the correct (U-shaped) way — and $\hat{\omega}^*$ lands in a plausible range for both instrument choices.

A Data-Side Check on $\hat{\lambda}$: Auxiliary Parameters $E[e]$, $\text{Med}[e]$

Does $\hat{\lambda}$ imply a sane detection probability at a typical evading firm?

- Approach: Schennach's own auxiliary-parameter device (her $E[U]$ worked example), applied to $e(M)$ directly, evaluated **post-hoc at the already-fitted** $(\hat{\theta}, \hat{\gamma})$. No re-optimization:

$$E[e(M)] - \mu_e = 0 \quad \text{vs.} \quad E[\mathbb{1}(e(M) \leq \theta_e) - \tfrac{1}{2}] = 0$$

Both evaluated via the same per-firm Metropolis chain the estimator itself uses
(`tilted_e_diag_A.cpp`, `Code/Rcpp/1200-stage2-elvis.cpp`)

— μ_e estimated from a running per-firm average (free) - θ_e , from the empirical median of pooled, thinned post-burn-in draws across all firms:

Instrument	$\hat{\mu}_e = E[e]$	$\hat{\theta}_e = \text{Med}[e]$	$q(\hat{\mu}_e) = \hat{\lambda}\hat{\mu}_e$	$q(\hat{\theta}_e)$
m_{it-1}^*	5,269	569	0.62%	0.067%
$\mathcal{W}_{it-2}$	3,858	184	0.23%	0.011%

- ▶ $\hat{\mu}_e \gg \hat{\theta}_e$ in both cases, a right-skewed evasion distribution: most firms evade modestly, a long right tail of aggressive evaders pulls the mean well above the median.
- ▶ All four implied probabilities inside $(0, 0.5)$ but are too low.

Implementation — Moment Set B

Adding nuisance parameters to A's 8 moments:

$$g(M, Z, \theta_{\text{smooth}}) = \begin{pmatrix} \text{Amoments} \\ \ln M_j - \theta_{m,j} \\ \omega(M, Z)_j - \theta_{\omega,j} \\ \omega(M, Z)_j^2 - \theta_{\omega^2,j} \\ \psi(M, Z)(\omega(M, Z)_j - \theta_{\omega,j}) \\ \psi(M, Z)(\omega(M, Z)_j^2 - \theta_{\omega^2,j}) \\ \psi(M, Z)(\ln M_j - \theta_{m,j}) \\ \psi(M, Z)\varepsilon(M, Z)\omega(M, Z) \end{pmatrix}$$


- ▶ Per-industry $\ln M$ targets ($\theta_{m,j}$, jointly estimated, $J = 29$ industries) and
- ▶ Omega-anchored heteroskedasticity block, $\theta_{\omega,j}$ and $\theta_{\omega^2,j}$, $j = 1, \dots, J$ are estimated from sample, before optimiation, so cheap, no work for the optimizer, except for the additional associated γ nuisance parameters (one per new moment row) that the optimizer must handle.

- ▶ **Same objective, sampler, two-phase λ search** as A, no methodology change, only the moment row list and parameter count grow.
- ▶ **Needed two fixes A didn't**, both scoped to B only (A's already-validated results untouched):
 - ▶ *Optimizer settings*: at A's defaults (`xtol_rel=1e-6`, `maxeval=1000`, no time cap), **every single B point hit the `maxeval` cap** — not real convergence. Loosened `xtol_rel` to `1e-4` (the objective has real Monte Carlo noise from the finite Metropolis chain; resolving parameter changes to 1-in-a-million is tighter than the noise floor supports at this dimension), raised `maxeval` to 5000, and added a `maxtime=600s` safety net per BOBYQA call (not previously set at all).
 - ▶ *Speed*: B's `cov(Ghat)` step was ~12% of total per-call time (measured, not assumed — an earlier “too small to matter” read came from a synthetic benchmark at the wrong scale). Replaced with Apple's Accelerate/CBLAS (`cblas_dsyrc`), fused into the same C++ call that builds the moment matrix so it never crosses back to R — 157.6ms $\rightarrow$ 138.25ms per call (~12.3% faster), verified bit-close (relative difference $\sim 3 \times 10^{-13}$) before being trusted.
- ▶ **Result**: overnight run (real scale, `n_burn=500`, `n_keep=1000`), **31 of 32 lambda points across both instruments converged for real** (`conv=1`, via `xtol_rel`). Only one non-final quantile-phase point hit the `maxtime` cap.

$\hat{L}_n(\lambda)$ — Moment Set B

Stage-2 ELVIS: CUE objective across the lambda grid

Moment set B, n_burn=500, n_keep=1000, corner_mode=include_zero



Reading the picture: both instruments' minimum sits at the smallest tested λ , rising away from it.

Obj fun value dropping more rapidly than A's and reaching lower levels, but the minimum is still off the left edge of the grid — a further extension is needed before either $\hat{\lambda}$ can be treated as final.

Headline Estimates — Moment Set B

Instru- ment	$\hat{\lambda}$	$\hat{L}_n$	$\hat{\delta}_0$	$\hat{\delta}_1$	$\hat{\delta}_2$	$\hat{\omega}^*$	$\max \gamma $	conv
m_{it-1}^*	5.88×10^{-7}	0.00113	-1.236	1.236	0.183	3.38	2.43	16/16
$\mathcal{W}_{it-2}$	5.88×10^{-7}	0.00553	-0.515	0.890	0.064	6.97	4.83	15/16

- ▶ m_{it-1}^* **agrees closely with A** ($\hat{\omega}^* = 3.38$ vs. A's 3.34 — near-exact).
- ▶ $\mathcal{W}_{it-2}$ **moved substantially**: A gave $\hat{\omega}^* = 2.35$, $\hat{\delta}_2 = 0.396$; B gives $\hat{\omega}^* = 6.97$, $\hat{\delta}_2 = 0.064$. Flatter cost curvature. Plausible (B carries real extra information, the per-industry $\ln M$ moments), but a large jump.

Interpreting $\hat{\omega}^*$ and $\hat{\lambda}$ Against the Data

Two more auxiliary-parameter checks, same as the $E[e]/\text{Med}[e]$ - where does $\hat{\omega}^* = \hat{\delta}_1/2\hat{\delta}_2$ actually sit in the tilted ω distribution — a typical firm, or a tail one? - extending the detection-probability check to e 's upper tail (P90/P95/P99), not just its mean/median.

	$\hat{\omega}^*$	percentile of $\hat{\omega}^*$	$\text{Med}(\omega)$	$\text{P75}(\omega)$	$\text{P95}(\omega)$
A, m_{it-1}^*	3.34	71.9%	2.79	3.42	4.20
B, m_{it-1}^*	3.38	71.0%	2.83	3.51	4.54
A, $\mathcal{W}_{it-2}$	2.35	29.9%	2.95	3.84	4.78
B, $\mathcal{W}_{it-2}$	6.97	99.9%	2.76	3.70	4.76

- ▶ $\mathcal{W}_{it-2}$: Under A, $\mathcal{W}_{it-2}$ puts $\hat{\omega}^*$ at the 29.9th percentile: cost is already rising for roughly 70% of firms. Not consistent with large/high productivity firms evade less: “cost rises past a low bar” than “cost rises specifically for the most productive.” Under B, the *same instrument* swings to the opposite extreme, 99.9th percentile, cost rising only for a sliver of 0.1% at the very top.
- ▶ m_{it-1}^* , stable at ~71st percentile in **both** moment sets. Cost rises for roughly the top quartile, consistent for both moment sets. More trustworthy instrument of the two.

	P90(e)	P95(e)	P99(e)	$\hat{\lambda} \cdot \text{P90}$	$\hat{\lambda} \cdot \text{P95}$	$\hat{\lambda} \cdot \text{P99}$
A, m_{it-1}^*	10,392	21,297	84,430	1.22%	2.50%	9.92%
B, m_{it-1}^*	6,010	13,470	62,740	0.35%	0.79%	3.69%
A, $\mathcal{W}_{it-2}$	7,720	15,666	57,918	0.45%	0.92%	3.40%
B, $\mathcal{W}_{it-2}$	19,829	42,873	212,589	1.17%	2.52%	12.49%

Even at the 99th percentile of (tilted) evasion, detection probability stays well under the 50% FOC ceiling in all four cases — no sign of $\hat{\lambda}$ pushing anyone near the boundary, aggressive-evader tail included.

Notes for Discussion

- ▶ **Three of the four fits share the same unresolved edge; one is a different situation.** A's $\mathcal{W}_{it-2}$, B's m_{it-1}^* , and B's $\mathcal{W}_{it-2}$ all bottom out at $\lambda = 5.877 \times 10^{-7}$, with $\hat{L}_n$ still falling as the refinement — which genuinely extended past the original quantile grid's own edge — approaches it: real evidence the true optimum lies further down for these three. A's m_{it-1}^* is different: its quantile-phase minimum landed on an *interior* point ($\lambda = 2.351 \times 10^{-6}$, correctly triggering `bracket_local_min`'s interior-neighbor branch, not the edge-extension one), so its refinement window (1.175×10^{-6} to 5.877×10^{-6}) never tested anything smaller than the quantile grid already had. Its own minimum ($\hat{\lambda} = 1.175 \times 10^{-6}$) sits at that window's edge, but nothing below it has ever been tried — an untested gap, not the same kind of evidence. Either way, extending the grid further down is the right next step for all four — just for a different reason in m_{it-1}^* 's case, and this is still the one blocking item before any of today's numbers should be called final.

- ▶ $\mathcal{W}_{it-2}$ is unstable across moment sets in a way that matters economically, not just numerically. $\hat{\omega}^*$'s percentile swings from 29.9% under

Estimating Tax Evasion by Deconvolution

Objective

- ▶ Empirical exercise to motivate structural model

What can be deconvolved and what can't

Convolution (Meister, 2009)

The density of the sum of two **independent** random variables is equal to the **convolution** of the densities of both addends; hence

$$f_V = \int f_\varepsilon(V - u)f_u(u)du$$

Under $\varepsilon \perp\!\!\!\perp M$, e and $e \perp\!\!\!\perp M$, with known f_ε and $u = \ln\left(\frac{M+e}{M}\right)$. Also $M^* = M + e$, and $V = u + \varepsilon$.

Then, I can deconvolve:

► f_u : $f_V = \int f_\varepsilon(V - u)f_u(u)du$

I don't need f_e , actually

What deconvolution recovers is the **overreporting log-ratio**,

$$u_{it} \equiv \ln \left(\frac{M_{it}^*}{M_{it}} \right) = \ln \left(1 + \frac{e_{it}}{M_{it}} \right),$$

- ▶ the *share of overreported over true inputs*
- ▶ the same object discussed previously, now given an explicit structural interpretation
- ▶ Since it's a ratio, it's adjusted by the firm's true materials level and hence informative about the intensity of tax evasion
- ▶ Had I recovered e , I would have to normalize somehow to compare firms of different levels of material utilization

Getting rid of the log

If I'm interested in only the density of the ratio, I can use a **density transformation**

Theorem 1.7.1 (Hogg et al. 2019): Let X be a continuous random variable with pdf $f_X(x)$ and support $\mathcal{S}_X$. Let $Y = g(X)$ where $g(x)$ is a one-to-one *[monotonic] differentiable* function on the support of X , $\mathcal{S}_X$. Denote the inverse of g by $x = g^{-1}(y)$ and $dx/dy = d[g^{-1}(y)]/dy$. Then, the pdf of Y is given by

$$f_Y(y) = f_X(g^{-1}(y)) \left| \frac{dx}{dy} \right|, \text{ for } y \in \mathcal{S}_Y$$

where the support of Y is the set $\mathcal{S}_Y = \{y = g(x) : x \in \mathcal{S}_X\}$

Getting density of the ratio

Therefore, with $X = u$ and $g^{-1}(y) = \ln(1 + y)$

$$f_{\frac{e}{M}}(y) = f_u(\ln(1 + y)) \left| \frac{1}{1 + y} \right| \quad (37)$$

Informal proof

If $U = g(R)$ for a strictly monotonic differentiable g , then

$$f_R(r) = f_U(g(r))|g'(r)|$$

Proof (via CDFs). Take g strictly increasing (your case). Then

$$F_R(r) = P(R \leq r) = P(g(R) \leq g(r)) = P(U \leq g(r)) = F_U(g(r)),$$

using that g increasing preserves the inequality direction.

Differentiate both sides in r :

$$f_R(r) = \frac{d}{dr}F_U(g(r)) = f_U(g(r))g'(r).$$

If g were decreasing instead, the inequality flips ($P(U \geq g(r)) = 1 - F_U(g(r))$), and differentiating introduces a minus sign that exactly cancels the sign of $g'(r)$, hence $|g'(r)|$ of the general formula.

Everything goes through unchanged conditional on $Z = z$ (just work with $F_{R|Z}(\cdot | z)$, $F_{u|Z}(\cdot | z)$ throughout).

Discussion

Deconvolution estimates the unconditional f_u because of how we set it up

- ▶ From f_u , we can learn the probability of a firm of overreporting u ,
 $P(U < u) = \int_0^u f_u(t)dt$, or
 - ▶ the expected overreporting $\mu_u = \int u f_u(u)du$
- ▶ This is the unconditional probability, the same for all firms regardless of their characteristics
- ▶ We might be interested in the *probability of overreporting conditional* on the firm size

What if we are interested in the conditional density?

$$\begin{aligned} f_{V|Z}(V | Z) &= \frac{f_{V,Z}(V, Z)}{f_Z(Z)} \\ &= \int f_{V|u,Z}(V | u, Z) f_{u|Z}(u | Z) du \\ &= \int f_{\varepsilon}(V - u) f_{u|Z}(u | Z) du \end{aligned}$$

using

- ▶ $f_{V|u,Z}(V | u, Z) = f_{\varepsilon|u,Z}(V - u | u, Z) - V = u + \varepsilon$ conditional on u
- ▶ $\varepsilon \perp (u, Z)$ **jointly**, so $f_{\varepsilon|u,Z} = f_{\varepsilon}$

with $u \not\perp Z$ (otherwise conditioning on Z would be pointless).

Estimating the Conditional distribution

- ▶ **Discrete/bucketed** Z (e.g. quintiles of K_{it} or L_{it}): the same 1-D deconvolution separately within each bucket.
 - ▶ We've already looked into $E[u|Z]$
- ▶ **Continuous** Z : nonparametric deconvolution with a continuous covariate
 - ▶ Literature: $f_{u|V}(u|V)$ (Wang and Ye 2015), $E[Z|u]$ (Fan and Truong 1993), and $f_{Z|u}$ (Huang and Zhou 2020)
 - ▶ $f_{u|Z}$: feasible using characteristic functions and kernels
- ▶ Ideally, what I want is $f_{e|\omega}$, but that's what f_{ψ} is, the object I will obtain with the structural approach

Conditional Deconvolution — Backup

What About the Full Distribution?

- ▶ So far: $E[u_{ijt} \mid Z_{ijt}]$, and only at 10 decile bins of Z
- ▶ Natural next question: the whole conditional density $f_{u|Z}(u \mid z)$, for Z continuous
 - ▶ Is the size gradient in the *level* of evasion, its *spread*, or its *shape*?

Aside: How Deconvolution Works

We never observe u_{ijt} directly — only $V_{ijt} = u_{ijt} + \varepsilon_{ijt}$, a “blurred” version of it

- ▶ A histogram of V mixes the true shape of u with the noise ε — the density of a sum of independent things is the *convolution* of their densities, $f_V = f_u * f_\varepsilon$
- ▶ Convolutions are hard to undo directly, but they become simple *multiplication* after a Fourier transform (the characteristic function):

$$\phi_V(t) = \phi_u(t) \phi_\varepsilon(t) \quad \implies \quad \phi_u(t) = \frac{\phi_V(t)}{\phi_\varepsilon(t)}$$

- ▶ Estimate ϕ_V from the data, divide by the already-known ϕ_ε , transform back $\rightarrow \hat{f}_u$
- ▶ This is exactly the (unconditional) trick behind the first-stage results (Carroll and Hall 1988; Stefanski and Carroll 1990)

Conditional Version: Add a Local Window

Same trick, done separately for each value of Z :

- ▶ Instead of using *all* firms to build one blurred distribution, use only firms with Z_i close to z — a moving local window, weighted by a kernel K
- ▶ Un-blur *that* local distribution the same way as before

$$\hat{\phi}_{V|Z=z}(t) = \frac{\sum_i K\left(\frac{Z_i - z}{h}\right) e^{itV_i}}{\sum_i K\left(\frac{Z_i - z}{h}\right)} \implies \hat{f}_{u|Z=z}(u) = \frac{1}{2\pi} \int e^{-itu} \frac{\hat{\phi}_{V|Z=z}(t)}{\phi_\varepsilon(t)} \omega(t) dt$$

- ▶ $\omega(t)$: a damping weight, standard in any deconvolution, to keep the division numerically stable
- ▶ Valid because $\varepsilon \perp (u, Z)$: within any Z -window, the noise is still independent of u — the same assumption used throughout this section

Is This Standard?

- ▶ (Fan and Truong 1993), already cited on the “is it feasible?” slide, addresses the *mirror-image* problem: a noisy covariate and a clean outcome (see also (Huang and Zhou 2020) for the conditional-density version of that same direction)
- ▶ Ours runs the other way — clean Z , noisy u — which is why the conditional *mean* never needed correcting: $\varepsilon \perp Z \Rightarrow E[V \mid Z] = E[u \mid Z]$
- ▶ Only the *density* needs deconvolving, and the tool is just the classical deconvolution kernel (Carroll and Hall 1988; Stefanski and Carroll 1990), applied inside a local window — a natural combination, not (to our knowledge) a named result in the literature

Do We Need the Full Density?

- ▶ The conditional mean already carries the headline story: evasion rises then falls with size — going further only pays off if it answers a question the mean can't
- ▶ The one moment that would matter: the conditional **variance** $\text{Var}(u \mid Z = z)$
 - ▶ Does “large firms don't evade” mean (a) *every* large firm individually evades a little — tight distribution near zero, like the corporations — or (b) a mix of some firms still evading heavily and others near zero that cancel out in the mean?
 - ▶ (a) vs. (b) is a real audit-targeting question: under (a), size alone is a good targeting rule at the top; under (b) it isn't — some large firms are still hiding substantial evasion behind the average
- ▶ If pushed further: the full nonparametric conditional deconvolution (previous slides) is the general method; cumulant/method-of-moments matching is a cheap over-identification check on the variance estimate, not a substitute for it

The Geography of Tax Evasion

How Does a Firm's Location Affect its Tax Evasion?

Firms located in different regions might face different probabilities of detection and different tax rates, and thus different incentives to evade taxes.

- ▶ Tax rates:
 - ▶ Departments and municipalities in Colombia can set their own tax rates similar to sales tax, for example the Industry and Commerce Tax
 - ▶ Special tax regimes for firms located in certain regions, such as the Free Trade Zones (Zonas Francas) and the Special Economic Zones (Zonas Económicas Especiales)
 - ▶ Special ownership model in Atlantic coast
 - ▶ Some variation in tax rates is observed in the data

How Does a Firm's Location Affect its Tax Evasion?

- ▶ Probability of detection: Two hypothesis
 1. Firms located near the tax authority's offices might have access to the tax authority (lobbying, bribery, etc.) and thus face a lower probability of detection
 2. Firms located farther away from the tax authority's offices might face a lower probability of detection, because the auditing cost for the tax authority increases with distance, and thus they are less likely to be audited

Data

In the dataset from Colombia, I have two geographic variables:

- ▶ `section_country_code`: Section of the country (eg. Antioquia, Atlantico, Bogota D.E., etc.)
- ▶ `metro_area_code`: Metropolitan Area (eg. Bogota, Cali, Medellin, Barranquilla).

Metro Area Summary Statistics

Metro Area	N	Market Share	Avg. Sales Tax Rate	Med. Sales Tax Rate	Avg. Tax Evasion	Med. Tax Evasion
Bogotá D.C., Soacha	14345	35.3	7.4	8.9	4.2	10.3
Medellín, Valle de Aburrá	8147	19.9	7.6	9.2	4.6	10.7
Rest of Colombia	5895	14.9	4.6	2.0	0.7	10.2
Cali, Yumbo	4538	11.4	7.2	8.4	4.2	10.1
Barranquilla, Soledad	2896	7.3	5.6	5.9	2.5	13.3
Bucaramanga, Giron, Floridablanca	2311	5.3	5.7	6.0	5.9	13.8
Pereira, Santa Rosa de Cabal, Dosquebradas	1128	2.7	7.6	9.1	11.1	21.1
Manizales, Villamaria	639	1.6	7.0	8.1	3.2	7.2
Cartagena	606	1.6	4.6	2.7	11.6	20.0

Section of the Country Summary Statistics

Depto	N	Market Share	Avg. Sales Tax Rate	Med. Sales Tax Rate	Avg. Tax Evasion	Med. Tax Evasion
Bogotá D.C.	13954	34.3	7.3	8.9	4.3	10.3
Antioquia	8467	20.7	7.5	9.1	3.4	9.7
Valle del Cauca	5780	14.6	6.7	7.0	4.1	10.3
Atlántico	2922	7.3	5.6	5.8	2.6	13.2
Santander	2585	6.0	5.6	6.0	5.5	12.8
Cundinamarca	1282	3.5	5.8	6.0	1.9	12.8
Risaralda	1161	2.8	7.4	8.9	11.0	21.0
Tolima	763	1.9	4.0	0.0	10.1	21.2
Caldas	710	1.8	6.7	6.8	-0.8	6.6
Bolívar	635	1.7	4.4	2.1	9.8	19.7
Norte de Santander	617	1.4	5.2	6.0	-7.6	2.9
Boyacá	290	0.7	5.8	6.0	1.8	8.8
Quindío	288	0.7	5.2	6.0	6.9	18.1
Nariño	255	0.6	3.4	0.0	14.7	25.1
Huila	182	0.5	4.4	0.0	10.1	19.2
Cauca	155	0.4	3.7	0.0	-0.4	-1.8
Meta	119	0.3	2.8	0.0	-10.5	12.4
Córdoba	126	0.3	3.7	0.3	21.1	25.1
Cesar	72	0.2	1.8	0.0	-11.8	16.9
Magdalena	73	0.2	5.0	1.3	-33.7	-13.2
Sucre	54	0.1	4.5	4.4	-5.7	-11.4
Putumayo	9	0.0	8.4	6.0	22.8	7.1
Caquetá	5	0.0	5.9	6.0	50.7	70.1
Chocó	1	0.0	3.1	3.1	-36.0	-36.0

Prelim Results (Metro Area)

Dependent Variable:	$\mathcal{V}$
Model:	(1)
<i>Variables</i>	
Bogotá D.C., Soacha	0.0485*** (0.0132)
Cali, Yumbo	0.0706*** (0.0174)
Medellín, Valle de Aburrá	0.0440** (0.0141)
Manizales, Villamaria	0.0359 (0.0512)
Barranquilla, Soledad	0.0309 (0.0266)
Bucaramanga, Giron, Floridablanca	0.0654** (0.0226)
Pereira, Santa Rosa de Cabal, Dosquebradas	0.1391*** (0.0355)
Cartagena	0.1402* (0.0758)
Rest of Colombia	0.0244 (0.0197)
<i>Fit statistics</i>	
Observations	32,373
R ²	0.00280
Adjusted R ²	0.00255


Clustered (plant & year) standard-errors in parentheses
*Signif. Codes: ***: 0.01, **: 0.05, *: 0.1*

Prelim Results (Departamento)

Dependent Variable: Model:	ψ (1)
<i>Variables</i>	
Antioquia	0.0331** (0.0143)
Atlántico	0.0323 (0.0267)
Bogotá D.C.	0.0484*** (0.0133)
Bolívar	0.1116 (0.0743)
Boyacá	0.0498 (0.0822)
Caldas	-0.0275 (0.0574)
Cauca	0.1291 (0.0779)
Cesar	-0.3661 (0.3888)
Córdoba	0.2684** (0.0882)
Cundinamarca	0.0371 (0.0449)
Huila	0.1464 (0.0816)
Magdalena	-0.4084* (0.2023)
Meta	-0.1142 (0.1742)
Nariño	0.1441* (0.0655)
Norte de Santander	-0.0894* (0.0427)
Quindío	0.0801 (0.0608)
Risaralda	0.1396*** (0.0349)
Santander	0.0629** (0.0227)
Sucre	0.0642 (0.1590)
Tolima	0.1092** (0.0414)
Valle del Cauca	0.0718*** (0.0152)
<i>Fit statistics</i>	
Observations	32,373
R ²	0.00895
Adjusted R ²	0.00834

Clustered (plant & year) standard errors in parentheses
 Signif. Codes: ***: 0.01, **: 0.05, *: 0.1

On the Map



Next Steps

- ▶ Are there any patterns? Do they make sense?
- ▶ Can any of these regions be linked to Free Trade Zones or Special Economic Zones? Since when are these in place?
 - ▶ Zonas Francas: started in 1958. During this period, there were changes made in 1981 and 1985, and 1991. 1981-1984, firms in Zonas Francas were exempt from income and sales taxes, as well as from import duties. They were under strict vigilance by the tax authority and confined to specific locations. Since to be eligible, your products had to be exported, most traditional firms choose not to locate in ZFs. Mostly multinationals and exporting Maquilas took advantage of ZFs. There were ZFs in Barranquilla, Cartagena, Santa Marta, and Buenaventura.
 - ▶ Zonas Económicas Especiales: started in 2020

Appendix

Appendix

B) Value of Reported Ghost Purchases Over Firm Total Purchases, by Firm Size




Figure 1: Figure 1B from Carrillo et al. (2022) ↵

Robust Check for Size Heterogeneity

$$\mathcal{V}_{ijt} = \sum_{z \in Z} \sum_{h \in 1,2} \mu_{z,h} x_{z,ijt}^h + \eta_{ijt} \quad (38)$$

Table 1: Coefficient estimates . 95% Bootstrap CI with 250 replications. Stars denote Bootstrap p-value significance at (***) 0.001, (**) 0.05, and (*) 0.1 level.

	IV: Lagged Revenue	OLS: Lagged Revenue
Revenue	1.8012*** [1.5744, 2.0744]	
Revenue ²	-0.7594*** [-1.0265, -0.5149]	
Revenue _{t-1}		1.456*** [1.274, 1.6735]
Revenue _{t-1} ²		-0.5586*** [-0.784, -0.339]
Labour	-0.7149*** [-0.9244, -0.5631]	-0.4281*** [-0.6214, -0.2994]
Labour ²	0.1051 [-0.0561, 0.3035]	-0.09 [-0.2306, 0.1051]
Capital	-0.1124 [-0.277, 0.085]	-0.022 [-0.1776, 0.1737]
Capital ²	-0.2332** [-0.4548, -0.0406]	-0.2997*** [-0.5255, -0.1106]
Exports	-0.1127 [-0.7196, 0.3843]	-0.2759 [-0.8917, 0.192]
Exports ²	0.066 [-0.483, 0.7506]	0.2503 [-0.2569, 0.9386]
Imports	-1.1048*** [-1.4765, -0.7346]	-1.1253*** [-1.5047, -0.7651]
Imports ²	1.3084*** [0.8792, 1.7581]	1.3469*** [0.9235, 1.8055]
Const.	-0.1309*** [-0.1662, -0.0985]	-0.1158*** [-0.1492, -0.0826]

Table 2: Coefficient estimates. Percentile of size of different variables. 95% Bootstrap CI with 250 replications. Stars denote Bootstrap p-value significance at (***) 0.001, (**) 0.05, and (*) 0.1 level.


	IV: Lagged Revenue	OLS: Lagged Revenue
Revenue	1.1491*** [1.0863, 1.2497]	
Revenue _{t-1}		0.9948*** [0.9362, 1.0738]
Labour	-0.6258*** [-0.7001, -0.5599]	-0.5044*** [-0.5691, -0.4402]
Capital	-0.3082*** [-0.3679, -0.2278]	-0.2737*** [-0.3309, -0.1943]
Exports	-0.1004 [-0.5756, 0.283]	-0.2292 [-0.7139, 0.1385]
Imports	-0.7292*** [-0.9883, -0.5182]	-0.7386*** [-1.0034, -0.5334]

Why the Negative Sign for Some Small Firms?

- ▶ Smaller firms are more likely to be Proprietorships and less formal
- ▶ Proprietorships report their industrial activity as part of their individual income taxes. They need to keep track of their personal expenses and their industrial activity expenses
- ▶ Because they are more likely to be less formal than larger firms, they may lack a rigorous expense report system
- ▶ Thus, they might be consistently underreporting their expenses

Densities of $\mathcal{V}$ and ε

V and Epsilon Densities by Industry



Empirical Exercises using $\mathcal{V}$

Empirical exercises remain valid as long as changes in V (tax evasion proxy) are not explained by changes in M (or ε) but only by e or e/M

- ▶ Changes in the proxy are explained by changes to incentives to adjust overreporting not firms adjusting their optimal choice of true inputs.
- ▶ Fiscal Reform of 1983:
 - ▶ Alternative argument, changes in true inputs explain the increase in the overreporting share. How? Suppose e is unchanged. This implies that after the policy reform unincorporated firms decrease their level of inputs by either an increase in relative prices or productivity shocks. Productivity shocks are unrelated to tax rates, to this is unlikely. According to the FOC, taxes do not show in the optimal choice of firms. Besides, we see an increase after the reform, while the reform made only a one-time change. So this alternative argument does not hold.
 - ▶ Our argument is more transparent: an increase in the sales tax rate lead to an increase in the overreporting share because a higher tax rate increases the marginal benefit of overreporting.

Empirical Exercises using $\mathcal{V}$ (Continued)

- ▶ Tax Evasion heterogeneity along firm size:
 - ▶ Alternative argument, correlation between true inputs and size explain why the overreporting share increases along the size distribution and drops at the top. How? fix e . CD could not be capturing the relationship between labour (or capital) and materials. A non-linear relationship could explain the pattern in the data.
 - ▶ However, the pattern is consistent with evidence from other papers with administrative data. As a robustness check, I also use the lag and double lag labour and capital and the pattern persist.

References

- Allingham, Michael G., and Agnar Sandmo. 1972. "Income Tax Evasion: A Theoretical Analysis." *Journal of Public Economics* 1 (November): 323–38. [https://doi.org/10.1016/0047-2727\(72\)90010-2](https://doi.org/10.1016/0047-2727(72)90010-2).
- Carrillo, Paul, Dave Donaldson, Dina Pomeranz, and Monica Singhal. 2022. *Ghosting the Tax Authority: Fake Firms and Tax Fraud*. NBER. <http://www.nber.org/papers/w30242>.
- Carroll, Raymond J., and Peter Hall. 1988. "Optimal Rates of Convergence for Deconvolving a Density." *Journal of the American Statistical Association* 83 (404): 1184–86. <https://doi.org/10.1080/01621459.1988.10478718>.
- Fan, Jianqing, and Young K. Truong. 1993. "Nonparametric Regression with Errors in Variables." *The Annals of Statistics* 21 (4): 1900–1925. <https://doi.org/10.1214/aos/1176349402>.